

# Lab 6: MongoDB Querying

CSL303: Database Management Systems

## Objective

To gain hands-on experience with the MongoDB query language. This lab covers data manipulation (insertion, updates, deletion) and a wide range of data retrieval techniques using the find command and the aggregation framework.

## Setup & Data Model

This lab requires a running MongoDB instance. You will be working with a university database consisting of three collections: instructors, courses, and students.

### Step 1: Populate the Database

Connect to your MongoDB instance using the mongosh shell. Create a new database named university and populate the collections by running the following commands.

```
// Switch to or create the university database  
use university;
```

```
// Populate the 'instructors' collection  
db.instructors.insertMany([  
  { _id: 1, name: "Dr. Turing", department: "CS", rank: "Professor" },  
  { _id: 2, name: "Dr. Codd", department: "CS", rank: "Professor" },  
  { _id: 3, name: "Dr. Ohm", department: "EE", rank: "Associate Professor" },  
  { _id: 4, name: "Dr. Curie", department: "Physics", rank: "Professor" }  
]);
```

```
// Populate the 'courses' collection  
db.courses.insertMany([  
  { _id: "CSL101", title: "Intro to Programming", credits: 3, instructor_id: 1 },  
  { _id: "CSL303", title: "Databases", credits: 4, instructor_id: 2 },  
  { _id: "EEL201", title: "Digital Circuits", credits: 3, instructor_id: 3 },  
  { _id: "PHL100", title: "Intro to Physics", credits: 4, instructor_id: 4 },  
  { _id: "CSL211", title: "Data Structures", credits: 4, instructor_id: 1 }  
]);
```

```
// Populate the 'students' collection  
db.students.insertMany([  
  { _id: 101, name: "Alice", major: "CS", gpa: 3.9, courses_enrolled: [  
    { cid: "CSL101", grade: "A" }, { cid: "CSL303", grade: "A" } ] },  
  { _id: 102, name: "Bob", major: "EE", gpa: 3.2, courses_enrolled: [  
    { cid: "EEL201", grade: "B" } ] },  
  { _id: 103, name: "Charlie", major: "CS", gpa: 3.5, courses_enrolled: [  
    { cid: "CSL101", grade: "B" }, { cid: "CSL211", grade: "A" } ] },  
  { _id: 104, name: "David", major: "Physics", gpa: 2.8, courses_enrolled: [  
    { cid: "PHL100", grade: "C" }, { cid: "CSL101", grade: "C" } ] },  
  { _id: 105, name: "Eve", major: "CS", gpa: 4.0, courses_enrolled: [] }  
]);
```

## Instructions for Submission

You must submit a single JSON file named `lab4_submission.json`. The file should contain a single JSON array of objects. Each object must have two keys: `"question_number"` and `"query"`. The value for the `"query"` key should be your complete MongoDB command as a single string.

### Submission Template:

```
[
  {
    "question_number": 1,
    "query": "db.students.insertOne({ _id: 106, name: \"Frank\", major: \"EE\", gpa: 3.1, courses_enrolled: [] });"
  },
  {
    "question_number": 2,
    "query": "db.courses.find({ credits: 4 });"
  }
]
```

---

## Exercises

### Part 1: Insert Operations

1. Insert a single new student with the following details: `_id: 106`, `name: "Frank"`, `major: "EE"`, `gpa: 3.1`, and an empty `courses_enrolled` array.
2. Insert two new courses at once:
  - `_id: "MEL110"`, `title: "Thermodynamics"`, `credits: 3`, `instructor_id: 4`.
  - `_id: "BI0101"`, `title: "Intro to Biology"`, `credits: 3`, `instructor_id: 4`.

### Part 2: Find Queries

3. Find all students majoring in "CS".
4. Find all courses that are worth 4 credits, but only display their titles and instructor IDs. Do not show the `_id` field.
5. Find all students with a GPA greater than or equal to 3.5.
6. Find all instructors who are either a "Professor" or an "Associate Professor". Use the `$in` operator.
7. Find the names and majors of all students who are enrolled in the "Databases" course (CSL303).
8. Find the names of all students who have received a grade of "A" in any course.
9. Find all students who are enrolled in exactly two courses. Use the `$size` operator.
10. Find all courses that do not have an instructor assigned (assume this means the `instructor_id` field is missing or null). Use the `$exists` operator.
11. Find all students who are enrolled in both "CSL101" and "CSL211". Use the `$all` operator.

### Part 3: Update Operations

12. The university has decided to rename the "CS" major to "Computer Science". Update all students in the "CS" major to their new major name. Use `updateMany`.
13. Alice (`_id: 101`) has just completed another course. Add a new course object `{ cid: "CSL211", grade: "A" }` to her `courses_enrolled` array. Use the `$push` operator.

14. Increase the GPA of all students in the "EE" major by 0.2 points. Use the \$inc operator.
15. Change the grade for student Bob (\_id: 102) in course "EEL201" to "A". (Hint: Use the arrayFilters option with the update command).

#### **Part 4: Delete Operations**

16. The "Thermodynamics" course (MEL110) has been cancelled. Delete it from the courses collection.
17. A student, David (\_id: 104), has dropped out. Delete him from the students collection.
18. Alice (\_id: 101) has decided to drop "CSL101". Remove only that course from her courses\_enrolled array. Use the \$pull operator.

#### **Part 5: Aggregation Framework**

19. Calculate the average GPA for each major. The output should have two fields: \_id (the major) and avgGpa.
20. Count the number of courses taught by each instructor. The output should show the instructor's ID and the number of courses they teach, named courseCount.